

Craig Horton Advisory

The Transformation Brief

Issue 14 · Friday 10 July 2026 · brief.craighortonadvisory.com

This week

Microsoft is changing **the work**

4,800 roles are going. The operating model is the real story.

By Craig Horton, AI Transformation Advisor

Microsoft is changing the work

4,800 roles are going. The operating model is the real story.

This week Microsoft announced it is eliminating around 2.1% of its global workforce, about 4,800 people, as part of its company transformation.

Amy Coleman, Microsoft's Chief People Officer, said those roles are not being replaced by AI, and I believe that statement, while also believing the operating model is changing underneath it.

AI changes the volume of coordination, the decisions a manager sees first, the work that moves between teams, and the judgement leaders are expected to make when an agent is now part of the flow.

A headcount reduction without redesigned workflows is a cost action, because the same handoffs remain, the same unclear ownership remains, and the people left behind inherit more pressure with less capacity to resolve it.

A question came to mind, which work has genuinely changed, which decision has moved, who owns the outcome once an agent is involved, and where does human judgement still have authority to intervene, challenge or stop the work?

This is where many AI programmes are heading, the operating model is starting to catch up.


From my office

This week I put the same agentic HR initiative from Issue 13 back through the AI BVF.

The verdict remained Fix, leaning towards Stop, because the strategic intent was good and the workflow had a credible value case, while the gaps were still in governance and change enablement, no named owner for the agent's decisions, no defined override role, and incomplete evidence for a person affected by the outcome.

The AI BVF pre-deployment actions were clear, map each decision affecting a person, name the business owner, the control owner and the person with authority to override, design the intervention point before the agent acts, complete data lineage, model-risk documentation and the Data Protection Impact Assessment, then re-score when the evidence is strong.

The run modelled EUR 5m to EUR 9m of organisational drag across a ten-week window, and although this was illustrative, the point is real, delay costs money, while premature deployment costs trust.



On the radar

- Trust in AI agents has reached the standards agenda, with the International Telecommunication Union launching work this week on agent identity, accountability and meaningful human control, because agents are becoming operating actors and governance now has to move into the operating design.

ITU newsroom, 9 July 2026

- The AI estate is spreading faster than oversight, with DigiCert reporting AI-related security incidents or vulnerabilities in 78% of organisations surveyed, alongside incomplete inventory and traceability, and a board cannot govern agents it cannot see.

DigiCert, 7 July 2026

- Infrastructure is now part of the value case, with Google Cloud reporting that 83% of organisations need material infrastructure change to capture the agentic opportunity, because data access, integration, compute cost and control architecture decide what reaches scale.

Google Cloud, 2026 State of infrastructure in the agentic AI era, 8 July 2026

Against the grain

"We are not replacing people with AI" may be true at the level of a role, but I am seeing enterprises change the work around the role, who prepares the decision, who reviews it, what evidence remains, and who carries accountability when the output is wrong.

That is why the redesign has to come before the project.

The reframing question

Before you approve a headcount action attached to an AI programme, can you show the workflow that changed, the decisions that moved, and the named person who owns the outcome?

Once again, I appreciate your interest and time to read this brief, and as always, look forward to your feedback and comments. Enjoy your weekend. Craig

Sources

- Microsoft, the latest in our company transformation, 6 July 2026 (blogs.microsoft.com)
- ITU, Focus Group on Trust and Identity for Humans and Agentic AI, 9 July 2026 (itu.int)
- DigiCert, AI security research, 7 July 2026 (digicert.com)
- Google Cloud, 2026 State of infrastructure in the agentic AI era, 8 July 2026 (cloud.google.com)
- AI BVF regulatory diagnostic, run 9 July 2026, aibvf-mcp, github.com/Craig-Horton/ai-bvf